

University of Stuttgart  
Germany

# Investigating the evolution of fisher information for neural network dynamics

Marc Sauter

ICP



University of stuttgart

A thesis presented for the degree of

*B.Sc.*

2023

Copyright © 2023 by Marc Sauter  
All Rights Reserved

Algebra is like sheet music. The important thing  
isn't can you read music, its can you hear it.

— Cristopher Nolan

# Acknowledgements

I would like to express ...

# Declaration

I, Marc Sauter, declare that ...

---

*Signature*

---

*Date*

## **Abstract**

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

**Keywords**— Keyword1 - Keyword2 - Keyword3

# Table of Contents

<b>1</b>	<b>Machine Learning Basics</b>	<b>1</b>
1.1	Introduction to machine learning . . . . .	1
1.2	Neural networks . . . . .	2
1.2.1	Neurons . . . . .	2
1.2.2	Neural networks . . . . .	4
1.2.3	Mathematical view . . . . .	5
1.3	Training of neural networks . . . . .	6
1.3.1	Datasets . . . . .	7
1.3.2	Loss function . . . . .	8
1.3.3	Optimization . . . . .	10
1.3.4	Other optimizers . . . . .	12
1.4	Which assumptions are actually necessary . . . . .	12
<b>2</b>	<b>Fisher information</b>	<b>14</b>
2.1	Use in Statistics . . . . .	14
2.2	Fisher information as Riemannian metric . . . . .	17
2.2.1	Differentiable manifolds . . . . .	18
2.2.2	Tangent space . . . . .	20
2.2.3	Riemannian metric and Fisher Information . . . . .	22
<b>3</b>	<b>Neural Tangent Kernel</b>	<b>25</b>
3.1	Derivation from Gradient Flow . . . . .	25
3.1.1	What we call time . . . . .	25
3.1.2	Derivation of the NTK . . . . .	27
3.2	Interpretation . . . . .	28

<b>Appendix A</b>	<b>Appendix A: some mathematical proofs</b>	<b>33</b>
A.1	Proof of eq. (2.1.3) . . . . .	33
A.2	Proof of eq. (2.2.15) . . . . .	35



# List of Figures

1.1	This figure aids the explanation of the general operating principle of neurons in neural networks as they are used in this thesis. The parameters $\omega_i$ are denoted in orange, the input activations $a_i$ in blue and the activation $e$ with its corresponding activation function $f$ in violet. . . . .	3
1.2	This figure shows an example of a neural network. It consists of multiple neurons connected to each other. This specific network has 2 input values, one output value and a width of 4 neurons for a total of 3 layers. The names of those input and output values correspond to the use case of identifying if a 2d point lies within a region in 2d space. . . . .	4
1.3	Displayed in this figure are 4 examples of input values for the MNIST dataset. This dataset represents handwritten digits from 0 to 9. Each input value consists of a 28 by 28 grid of grayscale pixel values. . . . .	7
1.4	This figure shows a visual explanation for the Gradient Descent algorithm. The dot markers show where a parameter that starts at the cross marker might end up if the algorithm is applied. . . . .	11
2.1	This figure shows a picture of a Galton board. Taken from [10]. . .	15
2.2	This figure shows the probability distributions for a Galton boards with different drop-in positions. The slots where the ball can end up are labeled by the value of $x_i$ . . . . .	15
2.3	This figure illustrates the manifold of normal distributions. As coordinate system, $\mu$ and $\sigma$ are used. Every point in this manifold represents a probability distribution, as indicated by the arrow. . .	19

- 2.4 This figure contains an example of a tangent space of a manifold,  
which in this case is a 2d-surface. . . . . 20
- 3.1 This graph shows the effect of assuming gradient flow. The dashed  
line along with the circle markers show the positions during reg-  
ular GD with finite step sizes, the solid line shows the path of the  
parameters under gradient flow. . . . . 26

## List of Tables

# List of Abbreviations

<b>AI</b> .....	Artificial Intelligence
<b>ML</b> .....	Machine Learning
<b>ReLU</b> .....	Rectified Linear Unit
<b>NN</b> .....	Neural Network
<b>GD</b> .....	Gradient Descent
<b>SGD</b> .....	Stochastic Gradient Descent
<b>NTK</b> .....	Neural Tangent Kernel
<b>FI</b> .....	Fisher Information

# 1 | Machine Learning Basics

## 1.1 Introduction to machine learning

Over 70 years ago in October of 1950, at a time when computers weighed several tons, could only perform a few thousand operations per second and the pinnacle of machine intelligence were analogous robots that could follow light sources[1], Alan Turing published a paper in the journal of Nature discussing the question "Can machines think?"[2]. In there, he tries to tackle that question by instead proposing a game he calls the "imitation game". This game puts a human, whom we will call A, in a room where he can communicate via written messages with two different entities, one of which being a human called B, the other being a machine. A's goal is to determine from this simple communication alone which of the two entities is the human. The goal of both the machine and B is convincing A of themselves being the human.

The largest execution of such an experiment to date took place in early 2023 in the form of an online chat portal, where players had two minutes to talk to either another human or an Artificial Intelligence (AI) without knowing the type of their interlocutor. After more than two million participants had played the game for a total of more than 15 million conversations, only 68 % of the attempted classifications were correct guesses.

All of the advanced AI-bots used in this experiment were achieved using machine learning methods, if what those AIs do is "thinking" is still heavily debated to this day. The Oxford Learning Dictionary defines machine learning as "a type of artificial intelligence in which computers use large amounts of data to learn how to do tasks rather than being programmed to do them"[3]. How we can do this with neural networks and use them to learn complex tasks is explained in the following

sections.

## 1.2 Neural networks

### 1.2.1 Neurons

The term "**neural network**" (NN) is a reference to the workings of nervous systems of animals. These systems consist of a net of neurons, which are biological cells that are intricately connected to other neurons through structures called synapses. These connections carry electrical pulses between the different neurons that can excite them when they exceed certain thresholds. These thresholds vary from neuron to neuron and change over time. When this excitation occurs, new pulses in turn propagate from the excited neuron outwards to possibly excite other neurons. This interplay between excitation and transmission through the network of neurons may create what we perceive as thinking.

Mathematical analyses of these systems in attempts to eventually understand and replicate this thinking process have been done as early as the 1940's[4], the artificial neural networks that we use for machine learning today are modern attempts to replicate them partly.

The neurons, which are the building bricks of the artificial neural networks, are mathematical entities that take a fixed number of scalar values as inputs and convert them into a single output value. For a more visual explanation let's take a look at fig. 1.1. This is a case with two inputs. The big circle in the middle is representing the neuron itself. It takes as input the activation values  $a_i$ , multiplies them with their corresponding weights  $\omega_i$ , sums them up, adds a bias value  $b$  and then applies the activation function to get the resulting output value. The input activations  $a_i$  correspond to the electronic pulses in the nerve system, with the strength of the pulse coded into the value itself. No pulse in the biological system




Figure 1.1: This figure aids the explanation of the general operating principle of neurons in neural networks as they are used in this thesis. The parameters  $w_i$  are denoted in orange, the input activations  $a_i$  in blue and the activation  $e$  with its corresponding activation function  $f$  in violet.

would be represented by an activation of zero in the mathematical model. The weights  $w_i$  are a representation of how important single input values are for the activation of whole neuron. In the biological counterpart this might correspond to how thick or conductive the connections between the nerve cells are. Finally, the combination of bias  $b$  and activation function describes how high the sum of the input-weight-pairs has to be to activate the neuron and how the resulting value for the activation of the neuron changes for higher input activations. For example, a very simple output activation function would be a heaviside function ( $f(x) = 1$  for  $x \geq 0$  and  $f(x) = 0$  for  $x < 0$ ). The neuron then outputs 1 if the sum of the input-weight-pairs is bigger than the negative bias  $-b$ , which results in a positive value for  $e$  and 0 otherwise. One could say that the neuron can only be on or off. A more common activation function that will be used for the rest of this thesis is the ReLU function (Rectified Linear Unit). This function is defined as

$$f(x) = \begin{cases} x, & \text{if } x \geq 0 \\ 0, & \text{otherwise} \end{cases}. \quad (1.2.1)$$




Figure 1.2: This figure shows an example of a neural network. It consists of multiple neurons connected to each other. This specific network has 2 input values, one output value and a width of 4 neurons for a total of 3 layers. The names of those input and output values correspond to the use case of identifying if a 2d point lies within a region in 2d space.

Here the neuron gets activated as soon as the sum of the input-weight-pairs is higher than  $-b$  and the value of the activation increases linearly with  $e$  for those high sums.

### 1.2.2 Neural networks

A neural network can be built from these neurons, by connecting the outputs of certain neurons to the inputs of others. For example let's take a look at fig. 1.2. This neural network consists of three layers of four neurons each, takes two values as input and outputs one value at the end. It could for example be trained for detecting if a point on a 2d-grid is inside or outside of a given region. The input values would be the  $x$  and  $y$  coordinates of the point and the output value could represent the predicted probability of the point being inside the desired region. How exactly this training would work will be explained in later sections.

This network is only a very specific example of what a neural network might look like. In reality, there are various kinds of networks used to learn different tasks. For this thesis we will only talk about "fully connected" or "dense" neural net-



works. This means that every neuron in the first layer will receive every possible input value, and every neuron in later layers will receive the output of every neuron in the previous layer as an input. The activation function is the same for every neuron in this case, but the weights and biases vary throughout. How the output of the network gets handled may still differ through different use cases. Most of the time, the output is generated by regular neurons and gets fed into a softmax at the end. Details on the softmax function can be found in [5]. For our understanding it's enough to know that the softmax function converts the values of the output nodes into probabilities that add up to 1, where a higher value results in a higher corresponding probability. The structure of a neural network is generally called the "architecture" of the network, with the hidden layers or simply layers referring to the columns of neurons in between the input values and output neurons and the amount of neurons per layer referred to as the "width" of the network. This is also denoted in fig. 1.2.

### 1.2.3 Mathematical view

The previous explanations have been very visual and step by step in an attempt to make the topic more accessible. However, these concepts can be broken down to rather short mathematical expressions.

To start off we can define the inputs as  $a_i^{(0)}$ ,  $i = 1, \dots, n$ , with  $n$  being the amount of inputs. The weights of the first hidden layer can be denoted by  $\omega_i^{(1)}$ ,  $i = 1, \dots, n$ . Further we define  $\omega_{i,j}^{(k)}$  as the weight that connects the  $i$ -th neuron in the  $k$ -th layer with the  $j$ -th neuron in the  $(k - 1)$ -th layer. The values that  $i$  and  $j$  can assume depend on the widths of the respective layers,  $k$  can reach values between 1 and the amount of hidden layers plus 1, for the output layer. The bias of the  $i$ -th neuron in the  $k$ -th layer is denoted as  $b_i^{(k)}$ . Using this notation we can instantly write out

the output of the  $i$ -th neuron in the  $k + 1$ -th layer as

$$a_i^{(k+1)} = f \left( \sum_j \left( \omega_{i,j}^{(k+1)} a_j^{(k)} \right) + b_i^{(k+1)} \right). \quad (1.2.2)$$

To actually calculate the value, the  $a_j^{(k)}$  values have to be recursively replaced with their full calculation until one arrives at the input values of the network.

For simplicity reasons we will refer to the weights  $\omega_{i,j}^{(k)}$  and biases  $b_i^{(k)}$  together as "**parameters**" of the neural network. We will denote these collected in one ordered set as  $\theta = \{\theta_i\}_{i=1}^N$  if  $N$  is the amount of weights and the amount of biases added together. The mapping of the parameters onto  $\theta$  is not important as long as it is known so that one can calculate the output of a neural network if given  $\theta$  in the same way as when one is given the parameters. Another possible representation that we will use often is to write this set as a vector  $\theta \in \mathbb{R}^N$ . A short notation for the output of the neural network will be explained in the next section.

### 1.3 Training of neural networks

In the previous sections we discussed how neural networks are built up. Now we will give an example of how they can be trained and used to perform specific tasks. Specifically, we will look at how they can be trained on data sets in a process called **supervised learning**. For this we fix the architecture of our neural network, which includes the organisation of neurons, the input and output handling and the activation function. The things we can vary during training to make our network perform better are the parameters  $\theta$ .




Figure 1.3: Displayed in this figure are 4 examples of input values for the MNIST dataset. This dataset represents handwritten digits from 0 to 9. Each input value consists of a 28 by 28 grid of grayscale pixel values.

### 1.3.1 Datasets

First we need to assume that we have a given dataset containing input-output pairs that represent the task we want the neural network to perform. The neural network is supposed to learn which output is supposed to be generated from an input by evaluating the given inputs and desired outputs. We will call those desired outputs "**targets**" to distinct them from the actual output of the neural network during training. A good example of this would be the MNIST database, which was first used in 1994 in [6] as a modification of an earlier database, and has since become a popular entry-level classification task for machine learning. Examples of input data for this database are shown in fig. 1.3. This dataset consists of various handwritten digits from 0 to 9 represented by a 28 by 28 grid of grayscale pixel values. The values of the pixels in these grids act as the 784 input values for the neural network, the targets should be shaped according to the output of the neural network. For example one might use a scalar output of the neural network that should be equal to the value of the digit drawn in the input pixel grid. In this case the targets are scalar values from 0 to 9 corresponding to their input pictures. Another way would be to use 10 output values together with the previously mentioned softmax function (see section 1.2.2) to receive the values as probabilities for the different numbers as output. We would then change our targets to vectors

with 10 elements, where the entry corresponding to the handwritten digit in the input picture would be 1, every other value would be 0.

We will denote the data sets as mathematical sets of input-target pairs  $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$ , where  $\mathbf{x}_i$  are the input values and  $\mathbf{y}_i$  are the targets. Here we called the amount of training data  $N$ . The mathematical dimension of  $\mathbf{x}_i$  and  $\mathbf{y}_i$  can vary and are dependent on the network architecture, but we will assume that they are vectors of the real numbers  $\mathbf{x}_i \in \mathbb{R}^a$   $\mathbf{y}_i \in \mathbb{R}^b$ . Due to this property we also sometimes refer to them as "points". If our inputs values are organized in a different manner, for example the MNIST data being matrices of  $\mathbb{R}^{28 \times 28}$  we can rearrange these numbers into a vector of  $\mathbb{R}^{784}$  in any way we want. The only important thing is that we convert every data point into a vector in the same way.

Going further we will denote all the output values of the neural network for the input  $\mathbf{x}_i$  by  $f_\theta(\mathbf{x}_i)$ . This means that we mathematically describe the whole behavior of the neural network as a function

$$f : \mathbb{R}^a \times \mathbb{R}^p \rightarrow \mathbb{R}^b \quad (1.3.1)$$

with  $p$  being the amount of parameters in the network.

### 1.3.2 Loss function

Now we have defined what a neural network is and how the input data we want to train on is structured. In order to train our networks to behave according to our data set, we now need to introduce a way to measure how well our network is performing. Once we can evaluate the performance of our network, we can then introduce ways to optimize that performance.

This measure of the performance of a neural network is called the "**loss function**"  $\mathcal{L}$ . It is also sometimes referred to as the cost function in literature. For this study

we will only consider loss functions of the form

$$\mathcal{L}(\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N, \theta) = \frac{1}{N} \sum_{i=1}^N \ell_{\theta}(\mathbf{x}_i, \mathbf{y}_i). \quad (1.3.2)$$

As a reminder,  $\theta$  is a set containing the parameters of the neural network. This is the variable of the loss function that describes the neural network. How one calculates the output of the neural network using the inputs and  $\theta$  has to be intrinsically defined as well. Sometimes it's more convenient to also denote this dependency, which would make the loss function

$$\mathcal{L}(\{(f_{\theta}(\mathbf{x}_i), \mathbf{y}_i)\}_{i=1}^N) = \frac{1}{N} \sum_{i=1}^N \ell_{\theta}(f_{\theta}(\mathbf{x}_i), \mathbf{y}_i). \quad (1.3.3)$$

How exactly the loss function should be defined cannot be generally stated. The only thing certain is that the value of the loss function should generally be smaller the closer the output of the neural network is to the corresponding targets.

In most of our cases we use loss functions with

$$\ell_{\theta}(\mathbf{x}_i, \mathbf{y}_i) = d(f_{\theta}(\mathbf{x}_i), \mathbf{y}_i), \quad (1.3.4)$$

where  $d$  represents the distance between the output vector of the neural network and the target vector according to different metrics. One very simple example would be

$$\mathcal{L}(\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N, \theta) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N (f_{\theta}(\mathbf{x}_i)_j - y_{i,j})^2, \quad (1.3.5)$$

with  $f_{\theta}(\mathbf{x}_i)_j$  and  $y_{i,j}$  being the  $j$ -th component of the network output and target vectors. Further examples and more general loss functions can be seen in [7].

### 1.3.3 Optimization

To quickly recap, we have now defined what a neural network is and that we want a fixed network architecture during training whose parameters we can vary to optimize the performance. We also assume that we have a data set that we want our network to behave like and a loss function that measures how well the current setup of our network is performing. We will now talk about how we can change the parameters in an attempt to optimize performance, which is achieved by lowering the value of the loss function.

The simplest and most common algorithm to do this with is **Gradient Descent** (GD). Here, the rule for updating the parameters looks like

$$\theta' = \theta - \eta \nabla_{\theta} \mathcal{L} [\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N, \theta], \quad (1.3.6)$$

where  $\eta$  is a parameter of training called the "learning rate", which is a scalar value for GD but could also vary through training and even be a tensor for more advanced optimization methods.

The idea behind this is that the gradient with respect to  $\theta$  points in the direction of the steepest ascent of the loss when the parameters are varied. If we subtract this gradient from the parameters we change the loss in the opposite direction, which is the direction of steepest descent. To visualize this, let's examine a single parameter that we will call  $\theta_1$ . The update rule for this single parameter can be obtained from writing out the previous eq. (1.3.6) in its full vectorial component form and then picking out the line of  $\theta_1$ . It reads

$$\theta'_1 = \theta_1 - \eta \frac{\partial}{\partial \theta_1} \mathcal{L} [\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N, \theta]. \quad (1.3.7)$$




Figure 1.4: This figure shows a visual explanation for the Gradient Descent algorithm. The dot markers show where a parameter that starts at the cross marker might end up if the algorithm is applied.

To explain the concept of this algorithm more visually let's take a look at fig. 1.4. Here the cross markings symbolize the position of parameter  $\theta_1$  before the update step. The algorithm then calculates the derivative of the loss  $\frac{\partial}{\partial \theta_1} \mathcal{L}$ , which is the slope of the tangent. Now, to change the parameter, the slope is multiplied with  $-1$  times the learning rate  $\eta$  and added to  $\theta_1$ , which is represented by the arrows. These arrows are adjusted in length so that the distance they cover in the direction of  $\theta_1$  is  $\eta$  times the derivative. They point in the direction of the tangent to show their dependence on the slope. One can see that the red arrow on the right covers more distance in the direction of  $\theta_1$  than the green one on the left, which is a result of the corresponding slope being steeper. The circle markers then represent where the parameter values end up after the update step.

This visual example already illustrates some of the drawbacks of the GD algorithm. For example, the green starting parameter will end up in a local minimum that has a higher loss value than the local minimum that the red parameter value tends to. Also going further to the left would make the loss function smaller than

both of these local minima. GD just tends to find the closest local minimum instead of a global one. Another drawback results from the step size being finitely big, which makes the green parameter hop over the local minimum. Both of these problems can't be solved by simply optimizing the learning rate  $\eta$  in general.

### 1.3.4 Other optimizers

The GD calculates the gradient from the whole data set, although in most cases it is sufficient or even beneficial to compute the loss function and its gradient for a subset of the whole data set, which is usually called a batch or mini-batch. The extreme end of this would be to calculate the gradient and update the parameters for one data point each, which is then called **stochastic Gradient Descent**. GD and the variations are very simple and fast to compute algorithms, but they can only find local minima and can quite easily hop over the actual minima. That's why other algorithms have been developed that try to solve or mitigate some of the problems of GD. For the training, which will be performed later in this thesis, the so-called ADAM optimizer was used. The intricacies of this optimization method are beyond the scope of this work. We refer the interested reader to [8], where the details of this algorithm are discussed. Another optimization algorithm that can be used to overcome some of the problems of GD is the Natural Gradient Descent, which will be mentioned in later sections.

## 1.4 Which assumptions are actually necessary

In the previous sections we took a brief look at machine learning and neural networks. Although this barely scratches the surface of the methods that are used today, it's still more than what's needed for the upcoming chapters. The specific methods were given to explain how the NNs that will come up later in this thesis



were trained, but are unnecessary restrictions that don't need to be made for the mathematical considerations coming up.

For those it is sufficient to view systems  $f_\theta(\mathbf{x}_i)$  that depend on parameters  $\theta$ , accept input vectors  $\mathbf{x}_i$  of constant dimension and can be evaluated through the formalism of the loss function from equation eq. (1.3.2). For the discussions of the NTK, we can even disregard the assumption of the loss function splitting up into a sum of  $\ell$  functions. This means that the exact properties of the neural networks we looked at earlier can be ignored, allowing us to generalize the observations to many more network architectures or completely different systems than previously mentioned.

## 2 | Fisher information

### 2.1 Use in Statistics

To introduce the **Fisher Information** (FI), we will start off with how it is defined and used in statistics.

Let's define a statistical model  $f(x_i|\theta)$  that represents how a parameter  $\theta$  is related to the outcomes  $x_i$  of random variables  $X_i$ [9]. As an example of what this means, let's look at a Galton board. For readers that are not familiar with what Galton boards are, there is an example photograph in fig. 2.1. It's a famous mechanical model that visualizes binomial distributions, which are approximations of the normal distribution. If we place many balls at the top of the board and let them fall to the bottom, the amount of balls that end up in each cell are distributed according to the binomial distribution. In this case,  $x_i$  could assume the slot number which a ball can fall into. The  $i$  could label multiple throws into the board. To introduce a parameter that influences the distribution of balls, let's say one can throw from different spots above the Galton board which we now control with the value of  $\theta$ . For a known  $\theta$ , the resulting value of the statistical model represents the probability distribution  $f(x_i|\theta) = p_\theta(x_i)$  for the probability of the different  $x_i$  outcomes. A visual representation of the probability distributions for different  $\theta$  can be seen in fig. 2.2.

In general, the statistical models might be more complex, with  $\theta$  containing multiple parameters,  $x$  being an element of a mathematical space other than  $\mathbb{R}$  and the index  $i$  denoting various different experiments that all depend on the same parameter but have different possible outcomes and probability distributions.

What's of interest for the FI are cases where the parameters are not known before



Figure 2.1: This figure shows a picture of a Galton board. Taken from [10].




Figure 2.2: This figure shows the probability distributions for a Galton boards with different drop-in positions. The slots where the ball can end up are labeled by the value of  $x_i$ .

conducting the experiment, and have to be approximated by the different outcomes  $x_i$ .

Before we introduce the FI, let's look at an example from [9]. Let's consider a biased coin where we denote the probability for heads ( $x_i = 1$ ) with  $\theta$  and the probability for tails ( $x_i = 0$ ) with  $1 - \theta$ . We will now take a look at the outcome of  $n$  throws represented by the variable  $X^n$ . For example one observed result for  $X^5$  could be  $x^5 = (1, 1, 0, 1, 0)$ . Let's consider another variable  $Y$  whose observation is the sum of the total head throws  $y = \sum x^n$ . In our example case of  $x^5 = (1, 1, 0, 1, 0)$  this would result in a value of  $y = 3$ . The probability for this variable  $y$  is distributed according to the binomial distribution, which means  $f(y|\theta) = \binom{n}{y} \theta^y (1 - \theta)^{n-y}$ . Here the binomial coefficient  $\binom{n}{y}$  represents the different combinations that result in the same  $y$  value. This is needed because there are  $2^n$  different possibilities for  $x$  while there are only  $n$  different possibilities for  $y$ .

If we now fix the outcome of  $y$  and look at the conditional probability for the different  $x^n$  that could have resulted in that  $y$  value, we get  $p(x|y, \theta) = 1/\binom{n}{y}$ . With  $p(x|y, \theta)$  we denote the probability depending on  $x$  for fixed  $y$  and  $\theta$ . Although the probability of  $y$  and  $x$  both depend on  $\theta$ , the probability for  $x$  when  $y$  is fixed doesn't. After measuring  $y$  there is no information about  $\theta$  left in the measurement of  $x$ . This means that  $y$  is fully descriptive of or "sufficient for" the parameter  $\theta$ . Measuring  $y$  results in the same amount of information about the parameter  $\theta$  as measuring the whole observation  $x$ . To quantify how much information a certain function of outputs  $t(x)$ , which is  $y(x^n)$  in our example, contains about the parameters  $\theta$ , Fisher introduced the Fisher information.

The Fisher information is defined as

$$I_{X,ij}(\theta) = E_{x \in X} \left[ \frac{d}{d\theta_i} \log f(x|\theta) \cdot \frac{d}{d\theta_j} \log f(x|\theta) \right], \quad (2.1.1)$$

where we used the expectation  $E$

$$E_{x \in X} [A(x)] = \begin{cases} \sum_{x \in X} (A(x)p(x)) & \text{if } X \text{ is discrete,} \\ \int_{x \in X} A(x)p(x)dx & \text{if } X \text{ is continuous.} \end{cases} \quad (2.1.2)$$

Since the Fisher information is dependent of  $\theta$ , we can assume the value of  $\theta$  to be fixed during calculation, which makes  $f(x|\theta)$  equal to  $p_\theta(x)$ .

For  $n$  independent experiments  $X^n$ , where  $f(x^n|\theta) = \prod_{i=1}^n f(x_i|\theta)$ , one can split the FI into

$$I_{X^n, ij}(\theta) = \prod_{i=1}^n I_{X_i, ij}(\theta). \quad (2.1.3)$$

A proof of this can be found in section [A.1](#).

For our example the FI yields  $I_{X^n}(\theta) = I_Y(\theta) = n/(\theta(1-\theta))$ [9]. This means that there is as much information about the  $\theta$  contained in the measurement of  $Y$  as in the measurement of  $X^n$ , which coincides with  $Y$  being a sufficient measurement for  $\theta$ .

To conclude this chapter, the Fisher information is used in statistics to measure the amount of information one can gather about a parameter  $\theta$  by measuring the outcome of a probability distribution  $p_\theta(x_i)$ . It is defined in equation eq. [\(2.1.1\)](#).

## 2.2 Fisher information as Riemannian metric

Previously we talked about the FI in the context of statistics. You may wonder why we went over a statistical method to describe probabilities, when we previously only talked about machine learning and neural networks. The answer to this question is, that the Fisher information matrix defined as

$$I_{ij} = E_{(\mathbf{x}_i, \mathbf{y}_i) \in D} \left[ \frac{d}{d\theta_i} \ell(\mathbf{x}_i, \mathbf{y}_i) \cdot \frac{d}{d\theta_j} \ell(\mathbf{x}_i, \mathbf{y}_i) \right] \quad (2.2.1)$$

acts as the Riemannian metric describing the statistical manifold of the network regarding its loss. A brief introduction to this topic will be provided in this chapter. Keep in mind that this is mostly for intuition purposes and we will only cover a few important definitions. For more details, please refer to [13] where most of the following information is taken from. If you're not interested in mathematical details you can skip this chapter and go directly to **FILL SOMETHING IN HERE**.

### 2.2.1 Differentiable manifolds

To state the definition, a  $n$ -dimensional manifold  $S$  is a topological space so that for every point you can define a neighborhood around that point which is homeomorphic to an open subset of  $\mathbb{R}^n$ . A good example of this would be the surface of the earth, where locally viewed in the scale that we usually see things, the earth appears flat, but on a global scale the earth is obviously a sphere. This results for example in the shortest path between two points not being a straight line in maps of the world as a whole. Also the angles of a triangle don't sum up to  $180^\circ$  as they would in a subspace of  $\mathbb{R}^n$ . This results from conventional maps being subspaces of  $\mathbb{R}^2$ , although the earth is only homeomorphic to  $\mathbb{R}^2$  in smaller local scales. If one tries to map the whole sphere into a map without gaps, one has to map the coordinates in a way that makes the shortest lines curved for example.

Let's come back to the statistical models  $f(x|\theta)$  we talked about in the last section. We will treat the models considering fixed parameters as probability distributions  $p_\theta(x)$  in this context. If the probabilities are sufficiently smooth in  $\theta$ , which means that they are differentiable in  $\theta$  as often as needed for further considerations, one can view the family of probabilities as a  $n$ -dimensional manifold, where the  $n$  different  $\theta$  components play the role of the coordinate system of the manifold.




Figure 2.3: This figure illustrates the manifold of normal distributions. As coordinate system,  $\mu$  and  $\sigma$  are used. Every point in this manifold represents a probability distribution, as indicated by the arrow.

For example let's consider normal distributions

$$p(x|\mu, \sigma) = \frac{1}{\sqrt{(2\pi\sigma^2)}} e^{-(x-\mu)^2/(2\sigma^2)}, \quad (2.2.2)$$

where  $\theta = \{\theta_1, \theta_2\} = \{\mu, \sigma\}$ . We can now consider this family of distributions as a manifold, displayed in fig. 2.3. This is like a space, where every point in the space represents a distribution  $p(x|\theta)$ .

It might also be clear that the coordinate system of a manifold is definable in multiple different ways. Although it's always given in our use case, let's therefore denote that in general when we have coordinates  $\theta$  we also need a mapping  $\phi$ , which maps coordinates to points on a manifold. This means that by applying  $\phi(p)$  to a point  $p \in S$  the resulting vector in  $\mathbb{R}^n$  resembles the coordinates of that point. We can also apply the inverse of that mapping to a set of coordinates to the point in the manifold that's represented by those coordinates.




Figure 2.4: This figure contains an example of a tangent space of a manifold, which in this case is a 2d-surface.

### 2.2.2 Tangent space

The tangent space  $T_p$  of a manifold at point  $p$  is a vector space obtained by linearization of the manifold around  $p$ . For intuition purposes, let's take a look at the tangent plane of a 2d-surface in section 2.2.2. Here the tangent space is simply a plane that touches the surface in one point, with derivatives adjusted to match the surface at that point. For the general case of  $n$ -dimensional manifolds, it is obvious that tangent spaces aren't simply tangent planes of surfaces in every case, therefore let's introduce a way how to calculate a tangent space.

First we will define curves  $c(t)$  that are continuous mappings from an interval  $[a, b] \in \mathbb{R}$  into the manifold  $S$ . In the parametric representation, the curve is given by  $\theta(t)$ . Now we can define what a tangent vector is.

Imagine a smooth real function  $f(\theta) : S \rightarrow \mathbb{R}$ . We can now restrict this function to our predefined curve  $c$  by  $f \circ c : [a, b] \rightarrow \mathbb{R}$ . We'll denote this via  $f(\{\theta(t)\})$  in



the coordinate expression. The derivative  $Cf$  of this function is then given by

$$Cf = \frac{df \circ c}{dt} = \sum_{i=1}^n \frac{\partial f}{\partial \theta_i} \frac{d\theta_i(t)}{dt} = \underbrace{\left( \sum_{i=1}^n \frac{d\theta_i}{dt} \frac{\partial}{\partial \theta_i} \right)}_{\text{Operator } C} f. \quad (2.2.3)$$

Therefore we can associate a directional derivative operator  $C$  with each curve. The only dependence of this operator regarding the curve is  $\frac{d\theta_i}{dt}$ , where we might also clarify that this derivative depends itself on the point where it is calculated at. If the manifold is infinitely differentiable, the set of these mappings  $C$  at a fixed point on the manifold forms a  $n$ -dimensional vector space, called the "**tangent space**"  $T_p$  of that point  $p$ .

To make this more clear, let's look at the simplest basis for this vector space, which we'll call the "natural basis" of the tangent space. For this we'll consider curves  $c_1, c_2, \dots, c_n$  through a point  $p_0$ , where the curves are defined as

$$c_i(t) = \{\theta_1^0, \theta_2^0, \dots, \theta_i^0 + (t - t_0), \dots, \theta_n^0\}, \quad (2.2.4)$$

so that  $c_i(t_0) = \{\theta_1^0, \theta_2^0, \dots, \theta_n^0\} = p_0$  for every curve. The tangent vectors  $C_i$  then are simply the derivatives regarding their corresponding coordinate  $C_i f = \partial/\partial \theta_i f$ . We will denote this in short by  $C_i = \partial_i$ . The  $n$  vectors  $\partial_i$  are linearly independent and form the natural basis for the tangent space. This means that any tangent vector  $A$  can be represented by

$$A = \sum_{i=1}^n A_i \partial_i, \quad (2.2.5)$$

with components with respect to the natural basis  $A_i$ . If we are given a curve  $c(t)$  going through  $c(t_0)$  which derivative operator  $C$  in  $t_0$  is equivalent to the vector

$A$ , we can find the natural basis representation of  $A$  as

$$A_i = \left. \frac{d\theta_i}{dt} \right|_{t_0}. \quad (2.2.6)$$

Now let's take a look at the case of manifolds of statistical models. First we will define  $\ell$  as

$$\ell(x|\theta) = \log f(x|\theta), \quad (2.2.7)$$

and assume that for fixed  $\theta$ , the  $n$ -functions  $\partial_i \ell(x|\theta)$  are linearly independent. This means that we can construct a vector space by defining

$$T_\theta^{(1)} = \{A(x) | A(x) = A_i \partial_i \ell(x|\theta)\}. \quad (2.2.8)$$

We do this because there is a natural isomorphism between the tangent space  $T_\theta$  and this space  $T_\theta^{(1)}$  through

$$\partial_i \in T_\theta \leftrightarrow \partial_i \ell(x|\theta) \in T_\theta^{(1)}. \quad (2.2.9)$$

We will call  $T_\theta(1)$  the "**1-representation**" of our tangent space for the statistical models. We will now use this vector to define an inner product on the tangent space and it's 1-representation.

### 2.2.3 Riemannian metric and Fisher Information

When the inner product of the tangent spaces  $T_p$  is defined, the manifold is called a **Riemannian manifold**.

Let's first consider the inner product of the 1-representation space. Let  $A(x)$  and

$B(x)$  be 1-representations of  $A$  and  $B \in T_\theta$ . It is intuitive to define the inner product as

$$\langle A(x), B(x) \rangle = E_{x \in X} [A(x)B(x)], \quad (2.2.10)$$

with the expectation value  $E[\cdot]$ . Since the tangent space is isomorphic to its 1-representation, the inner product also translates via

$$\langle A, B \rangle = \langle A(x), B(x) \rangle. \quad (2.2.11)$$

This also means that we can calculate the inner product of the basis vectors as

$$\begin{aligned} g_{ij}(\theta) &:= \langle \partial_i, \partial_j \rangle = \langle \partial_i \ell(x|\theta), \partial_j \ell(x|\theta) \rangle \\ &= E_{x \in X} [\partial_i \ell(x|\theta), \partial_j \ell(x|\theta)]. \end{aligned} \quad (2.2.12)$$

The resulting object  $g_{ij}(\theta)$  is called the **Riemannian metric tensor** of the manifold. We can see that the Riemannian metric tensor for the statistical model is equivalent to the Fisher Information. It might be of interest to note here that we assume that  $\ell$  only depends explicitly on  $\theta$ . If we denote these as implicit dependencies, we have to replace the partial derivatives with absolute ones.

The inner product of two vectors can now be expressed with the metric tensor as

$$\langle A, B \rangle = \sum_{i,j} A_i B_j g_{ij}(\theta) \quad (2.2.13)$$

in the component form.

Using this representation of the inner product, we can define various things. For example, the length of a vector  $A$  is defined as  $|A|^2 = \sum_{i,j} A_i A_j g_{ij}$ , the orthogonality of two vectors when their inner product is zero, and the distance between

two points  $\theta^{(0)}$  and  $\theta^{(1)}$  along the curve  $c$  is defined by

$$s = \int_{t_0}^{t_1} \sum_{i,j} \sqrt{g_{ij} \frac{d\theta_i}{dt} \frac{d\theta_j}{dt}} dt. \quad (2.2.14)$$

This also introduces the concept of **Riemannian geodesics**. These are the curves that connect two points via the minimal distance between the two.

Another representation of the metric tensor or the FI is

$$g_{ij}(\theta) = -E_{x \in X} \left[ \partial_i \partial_j \ell(x|\theta) \right]. \quad (2.2.15)$$

A proof of this is denoted in section [A.2](#).

## 3 | Neural Tangent Kernel

### 3.1 Derivation from Gradient Flow

#### 3.1.1 What we call time

A simple and intuitive way to introduce the NTK is via "**Gradient Flow**", which is an assumption related to the GD algorithm from section 1.3.3. To quickly recap the update step from stochastic gradient descent, it is defined as

$$\theta' = \theta - \eta \nabla_{\theta} \mathcal{L} \left( \{ (f_{\theta}(\mathbf{x}_i), \mathbf{y}_i) \}_{i=1}^N \right). \quad (3.1.1)$$

The "flow" aspect arises when we start to ignore the discrete nature of these update steps and assume that the  $\theta$  parameters change continuously. A visual demonstration of how this changes the evolution of the parameters can be seen in fig. 3.1. Here, the dashed markers represent the evolution under regular GD, while the solid line represents the evolution for gradient flow. To do this, we first introduce a notion of "time" into our system. We try to visualize the optimization process as an evolution of our parameters  $\theta$  through this variable called time, which converts updating the parameters into moving further along the timeline of our parameters. We translate  $\theta \rightarrow \theta(t)$  and  $\theta' \rightarrow \theta(t + \Delta t)$ . This time in our system doesn't work exactly the same way as physical time, but since the process of calculating better parameters and changing them is always associated with the expenditure of physical time, it is intuitive to refer to our system's propagation variable as "time". Returning to the GD algorithm, since the learning rate  $\eta$  affects the size of our update step, we will refer to  $\eta$  as the amount of time it takes to update a parameter




Figure 3.1: This graph shows the effect of assuming gradient flow. The dashed line along with the circle markers show the positions during regular GD with finite step sizes, the solid line shows the path of the parameters under gradient flow.

$\theta' \rightarrow \theta(t + \eta)$ . The whole GD algorithm then becomes

$$\theta(t + \eta) = \theta(t) - \eta \nabla_{\theta(t)} \mathcal{L} \left( \{(f_{\theta(t)}(\mathbf{x}_i), \mathbf{y}_i)\}_{i=1}^N \right) \quad (3.1.2)$$

which we can rewrite to

$$\frac{\theta(t + \eta) - \theta(t)}{\eta} = -\nabla_{\theta(t)} \mathcal{L} \left( \{(f_{\theta(t)}(\mathbf{x}_i), \mathbf{y}_i)\}_{i=1}^N \right). \quad (3.1.3)$$

When observing the term on the left side, readers who are familiar with calculus might recognize that it looks similar to the definition of a derivative in time

$$\frac{d}{dt} \theta(t) = \lim_{\eta \rightarrow 0} \frac{\theta(t + \eta) - \theta(t)}{\eta}. \quad (3.1.4)$$

This means that for very small learning rates we can approximate the GD as

$$\frac{d}{dt} \theta(t) = -\nabla_{\theta(t)} \mathcal{L} \left( \{(f_{\theta(t)}(\mathbf{x}_i), \mathbf{y}_i)\}_{i=1}^N \right) \quad (3.1.5)$$

with the partial derivative of  $\theta$  along the assumed time variable.

For visual simplicity reasons, let's define the  $j$ -th component of the network output for the  $i$ -th input point  $f_{\theta(t)}(\mathbf{x}_i)_j$  as  $f_{ij}$  and assume Einstein summation for the rest of the chapter. This means that when an index is occurring twice, we don't denote a hidden summation over all possible values for this index (for example  $a_k b_k = \sum_k a_k b_k$ ).

Because it will be convenient later, we also spell out one component of the  $\nabla_{\theta}$  derivation of the loss function further by using the chain rule as

$$\begin{aligned} \frac{d\theta_k}{dt} &= -\frac{d}{d\theta_k} \mathcal{L}(\{(f_{\theta(t)}(\mathbf{x}_i), \mathbf{y}_i)\}_{i=1}^N) \\ &= -\frac{\partial \mathcal{L}}{\partial f_{ij}} \cdot \frac{df_{ij}}{d\theta_k}. \end{aligned} \quad (3.1.6)$$

### 3.1.2 Derivation of the NTK

This notion of time affects not only the parameters, but also everything that depends on them. For example, since the network output of a fixed architecture for a given input data point only depends on the parameters of the network, it can also be mathematically viewed as dependent on the time  $f_{\theta}(\mathbf{x}_i) \rightarrow f_{\theta(t)}(\mathbf{x}_i)$ . This means we can also calculate the derivative of one of the network outputs  $f_{\theta(t)}(\mathbf{x}_i)_j = f_{ij}$  to

$$\begin{aligned} \frac{d}{dt} f_{ij} &= \frac{\partial f_{ij}}{\partial \theta_k} \frac{d\theta_k}{dt} \\ &= \frac{\partial f_{ij}}{\partial \theta_k} \left( -\frac{\partial \mathcal{L}}{\partial f_{ij}} \frac{df_{ij}}{d\theta_k} \right) \\ &= -\underbrace{\frac{\partial f_{ij}}{\partial \theta_k} \frac{\partial f_{lm}}{\partial \theta_k}}_{=: \Lambda_{iljm}} \frac{\partial \mathcal{L}}{\partial f_{lm}}. \end{aligned} \quad (3.1.7)$$

In the last line, we used the fact that  $f_{ij}$  only depend explicitly on  $\theta$ , which means that the total and partial derivative are interchangeable. The rank 4 hypermatrix  $\Lambda$  is what we call the Neural Tangent Kernel for GD. We sorted the indices of this matrix so that the first two refer to the input points of  $f$  and the last two refer to the components of the output dimensions of the neural network. Note that this NTK is derived directly from the update algorithm of the GD for infinitely small  $\eta$ , the equations above don't hold for other optimization systems.

Another way to derive the NTK using an approximation of  $\Delta\mathcal{L}$  for small  $\eta$  can be seen around page 196 of [11]. The NTK derived there also works for tensorial gradient descent with a learning rate tensor  $\eta_{ij}$ .

## 3.2 Interpretation

Through eq. (3.1.7) the NTK is defined as

$$\Lambda_{ij\alpha\beta} = \nabla_{\theta} f_{\theta}(\mathbf{x}_i)_{\alpha} \cdot \nabla_{\theta} f_{\theta}(\mathbf{x}_j)_{\beta}. \quad (3.2.1)$$

For now, let's assume that the neural network has only one output  $f_{\theta}(\mathbf{x}_i)_{\alpha} \rightarrow f_{\theta}(\mathbf{x}_i)$ , which gets rid of the  $\alpha$  and  $\beta$  indices for us and makes the NTK a regular matrix. We can use this matrix to calculate the time derivative of the neural network output similar to eq. (3.1.7) by

$$\frac{d}{dt} f_{\theta(t)}(\mathbf{x}_i) = \sum_j \Lambda_{ij} \left( -\frac{\partial \mathcal{L}}{\partial f_{\theta(t)}(\mathbf{x}_j)} \right). \quad (3.2.2)$$

This means that the evolution of the network output for input  $\mathbf{x}_i$  is influenced by the outputs for other input values  $\mathbf{x}_j$  through the NTK. We can investigate this further by taking a look at the definition of the NTK above in eq. (3.2.1).



A mathematical kernel  $K$  is defined as a function

$$K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j), \quad (3.2.3)$$

with another so-called "feature map" function  $\phi$  that maps the points  $\mathbf{x}_i$  into a higher dimensional inner product space called the "feature space". The kernel  $K$  assigns a scalar value to the two points by comparing their "features" via a scalar product in the feature space. For more specific information on Kernels and how they are used, please refer to [12].

In our case the feature map is  $\nabla_\theta$  which maps our scalar network output onto a vector that contains the derivatives of the network output with respect to all of the various parameters. This is our feature space, because what we're interested in is how moving in  $\theta$  space changes the output values of our network. If we compare those two mappings we get a value that measures how closely the direction that increases  $f_\theta(\mathbf{x}_i)$  most effectively matches with the direction that increases  $f_\theta(\mathbf{x}_j)$  most effectively (which are the directions the gradients point to). For example,  $\Lambda_{ij} = 0$  means that varying  $\theta$  in the direction that changes  $f_\theta(\mathbf{x}_i)$  most effectively results in 0 change for  $f_\theta(\mathbf{x}_j)$ .

Coming back to eq. (3.2.2) the right side is the negative loss derivative with respect to  $f_\theta(\mathbf{x}_j)$ . Since we update the parameters in a way that minimizes the loss most effectively in gradient flow, the negative loss derivative with respect to  $f_\theta(\mathbf{x}_j)$  describes how beneficial increasing  $f_\theta(\mathbf{x}_j)$  is for our system. If we now multiply this with the NTK, which is a kernel that describes how similar  $f_\theta(\mathbf{x}_i)$  and  $f_\theta(\mathbf{x}_j)$  behave when changing theta, and sum everything up, we get the change of the function value  $f_\theta(\mathbf{x}_i)$  as result.

In  $\theta$  space, we can directly relate the evolution of  $\theta$  to its negative loss derivative,

because the parameters themselves are what we update with our algorithm. For the derivative of the output values we need the NTK as well, because we don't change the output values directly. If the loss derivative with respect to an output value tells the system that it *should increase this output value*, the system *changes the parameters* in a way that increases this output value. The difference to the derivative of  $\theta$  is that the change of parameters now doesn't reflect in just one output, but in every other output value as well. That's where the NTK arises in the equation. The NTK is, in a way, a translation of the GD into the space of outputs of the neural network.

All of the explanations above apply to the 4 dimensional hypermatrix form of the NTK for multidimensional neural network output as well, which can be seen when comparing eq. (3.2.2) to eq. (3.1.7).

# Literature

- [1] Esther Inglis-Arkell. Mar. 7, 2012. URL: <https://gizmodo.com/the-very-first-robot-brains-were-made-of-old-alarm-cl-5890771> (visited on 08/23/2023).
- [2] A. M. TURING. “I.COMPUTING MACHINERY AND INTELLIGENCE”. In: *Mind* LIX.236 (Oct. 1950), pp. 433–460. ISSN: 0026-4423. DOI: [10.1093/mind/LIX.236.433](https://doi.org/10.1093/mind/LIX.236.433). eprint: <https://academic.oup.com/mind/article-pdf/LIX/236/433/30123314/lix-236-433.pdf>. URL: <https://doi.org/10.1093/mind/LIX.236.433>.
- [3] URL: <https://www.oxfordlearnersdictionaries.com/us/definition/english/machine-learning> (visited on 08/23/2023).
- [4] Warren S. McCulloch and Walter Pitts. “A logical calculus of the ideas immanent in nervous activity”. In: *Bulletin of Mathematical Biophysics* 5 (1943), pp. 115–133. DOI: <https://doi.org/10.1007/BF02478259>.
- [5] Bolin Gao and Lacra Pavel. “On the Properties of the Softmax Function with Application in Game Theory and Reinforcement Learning”. In: (2018). arXiv: [1704.00805 \[math.OC\]](https://arxiv.org/abs/1704.00805).
- [6] L. Bottou et al. “Comparison of classifier methods: a case study in handwritten digit recognition”. In: *Proceedings of the 12th IAPR International Conference on Pattern Recognition, Vol. 3 - Conference C: Signal Processing (Cat. No.94CH3440-5)*. Vol. 2. 1994, 77–82 vol.2. DOI: [10.1109/ICPR.1994.576879](https://doi.org/10.1109/ICPR.1994.576879).
- [7] Qi Wang et al. “A Comprehensive Survey of Loss Functions in Machine Learning”. In: *Annals of Data Science* 9 (2 2022), pp. 2198–5812. DOI: [10.1007/s40745-020-00253-5](https://doi.org/10.1007/s40745-020-00253-5).

- [8] Diederik P. Kingma and Jimmy Ba. “Adam: A Method for Stochastic Optimization”. In: (2017). arXiv: [1412.6980](https://arxiv.org/abs/1412.6980) [cs.LG].
- [9] Alexander Ly et al. “A Tutorial on Fisher information”. In: *Journal of Mathematical Psychology* 80 (2017), pp. 40–55. ISSN: 0022-2496. DOI: <https://doi.org/10.1016/j.jmp.2017.05.006>. URL: <https://www.sciencedirect.com/science/article/pii/S0022249617301396>.
- [10] *atematica (IME/USP)/Rodrigo Tetsuo Argenton*. URL: [https://commons.wikimedia.org/wiki/File:Galton\\_box.jpg#/media/Datei:Galton\\_box.jpg](https://commons.wikimedia.org/wiki/File:Galton_box.jpg#/media/Datei:Galton_box.jpg) (visited on 08/24/2023).
- [11] Daniel A. Roberts, Sho Yaida, and Boris Hanin. *The Principles of Deep Learning Theory*. Cambridge University Press, May 2022. DOI: [10.1017/9781009023405](https://doi.org/10.1017/9781009023405). URL: <https://doi.org/10.1017/9781009023405>.
- [12] Colin Campbell. “An introduction to kernel methods”. In: *Studies in Fuzziness and Soft Computing* 66 (2001), pp. 155–192.
- [13] Shun-ichi Amari. *Differential-Geometrical Methods in Statistics*. 1st ed. ISBN: 978-0-387-96056-2. DOI: <https://doi.org/10.1007/978-1-4612-5056-2>.
- [14] Daniel Jannai et al. *Human or Not? A Gamified Approach to the Turing Test*. 2023. arXiv: [2305.20010](https://arxiv.org/abs/2305.20010) [cs.AI].
- [15] S. Amari and S.C. Douglas. “Why natural gradient?” In: 2 (1998), 1213–1216 vol.2. DOI: [10.1109/ICASSP.1998.675489](https://doi.org/10.1109/ICASSP.1998.675489).
- [16] Enzo Grossi and Massimo Buscema. “Introduction to artificial neural networks”. In: *European Journal of Gastroenterology and Hepatology* 19.12 (Dec. 2007), pp. 1046–1054. DOI: [10.1097/MEG.0b013e3282f198a0](https://doi.org/10.1097/MEG.0b013e3282f198a0).

## A | Appendix A: some mathematical proofs

### A.1 Proof of eq. (2.1.3)

For all variables  $X^n$  that satisfy

$$f(x^n|\theta) = \prod_{\alpha=1}^n f(x_\alpha|\theta) \quad (\text{A.1.1})$$

we can prove

$$I_{X^n,ij} = \sum_{\alpha} I_{X_\alpha,ij} \quad (\text{A.1.2})$$

by

$$\begin{aligned}
I_{X^n,ij} &= E_{x^n \in X^n} \left\{ \frac{d}{d\theta_i} \log f(x^n|\theta) \frac{d}{d\theta_j} \log f(x^n|\theta) \right\} \\
&= \sum_{x^n \in X^n} \left\{ \left[ \frac{d}{d\theta_i} f(x^n|\theta) \right] \left[ \frac{d}{d\theta_j} f(x^n|\theta) \right] [f(x^n|\theta)]^{-1} \right\} \\
&\stackrel{\text{A.1.1}}{=} \sum_{x^n \in X^n} \left\{ \frac{d}{d\theta_i} \left[ \prod_{\alpha} f(x_{\alpha}|\theta) \right] \frac{d}{d\theta_j} \left[ \prod_{\beta} f(x_{\beta}|\theta) \right] \left[ \prod_{\gamma} f(x_{\gamma}|\theta) \right]^{-1} \right\} \\
&= \sum_{x^n \in X^n} \left\{ \left( \prod_{\alpha} f(x_{\alpha}|\theta) \right) \left( \sum_{\alpha} \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right) \right. \\
&\quad \left. \left( \prod_{\beta} f(x_{\beta}|\theta) \right) \left( \sum_{\beta} \frac{d}{d\theta_i} \log f(x_{\beta}|\theta) \right) \left[ \prod_{\gamma} f(x_{\gamma}|\theta) \right]^{-1} \right\} \\
&= E_{x^n \in X^n} \left\{ \left[ \sum_{\alpha} \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right] \left[ \sum_{\beta} \frac{d}{d\theta_i} \log f(x_{\beta}|\theta) \right] \right\} \\
&\stackrel{(*)}{=} E_{x^n \in X^n} \left\{ \sum_{\alpha} \left[ \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right] \left[ \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right] \right\} \\
&= \sum_{\alpha} E_{x^n \in X^n} \left\{ \left[ \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right] \left[ \frac{d}{d\theta_i} \log f(x_{\alpha}|\theta) \right] \right\} \\
&= \sum_{\alpha} I_{X_{\alpha},ij}.
\end{aligned}
\tag{A.1.3}$$

The equality at  $(*)$  holds because for all  $\alpha \neq \beta$

$$\begin{aligned}
& E_{x^n \in X^n} \left\{ \left[ \frac{d}{d\theta_i} \log f(x_\alpha|\theta) \right] \left[ \frac{d}{d\theta_i} \log f(x_\beta|\theta) \right] \right\} \\
& \propto \sum_{x_\alpha} \sum_{x_\beta} \left\{ \left[ \frac{d}{d\theta_i} f(x_\alpha|\theta) \right] \left[ \frac{d}{d\theta_i} f(x_\beta|\theta) \right] \right\} \\
& = \sum_{x_\alpha} \left\{ \left[ \frac{d}{d\theta_i} f(x_\alpha|\theta) \right] \sum_{x_\beta} \left[ \frac{d}{d\theta_i} f(x_\beta|\theta) \right] \right\} \\
& = \sum_{x_\alpha} \left\{ \left[ \frac{d}{d\theta_i} f(x_\alpha|\theta) \right] \sum_{x_\beta} \frac{d}{d\theta_i} [f(x_\beta|\theta)] \right\} \\
& = \sum_{x_\alpha} \left\{ \left[ \frac{d}{d\theta_i} f(x_\alpha|\theta) \right] \sum_{x_\beta} \frac{d}{d\theta_i} [1] \right\} \\
& = 0.
\end{aligned} \tag{A.1.4}$$

This prove still holds when using continuous variables instead of discrete ones.

## A.2 Proof of eq. (2.2.15)

Here we will prove eq. (2.2.15) which states

$$E_{x \in X} \left[ \partial_i \log p(x|\theta) \cdot \partial_j \log p(x|\theta) \right] = - E_{x \in X} \left[ \partial_i \partial_j \log p(x|\theta) \right]. \tag{A.2.1}$$

To prove this we will evaluate the right side by

$$\begin{aligned}
- E_{x \in X} [\partial_i \partial_j \log p(x|\theta)] &= - E_{x \in X} \{ \partial_i [\partial_j \log p(x|\theta)] \} \\
&= - E_{x \in X} \left\{ \partial_i \left[ \frac{\partial_j p(x|\theta)}{p(x|\theta)} \right] \right\} \\
&= - E_{x \in X} \left[ \frac{\partial_i \partial_j p(x|\theta)}{p(x|\theta)} - \frac{\partial_i p(x|\theta) \partial_j p(x|\theta)}{p(x|\theta)^2} \right] \\
&= - E_{x \in X} \left[ \frac{\partial_i \partial_j p(x|\theta)}{p(x|\theta)} - \partial_i \log p(x|\theta) \partial_j \log p(x|\theta) \right] \\
&= E_{x \in X} [\partial_i \log p(x|\theta) \cdot \partial_j \log p(x|\theta)],
\end{aligned} \tag{A.2.2}$$

Where the last equation holds because

$$\begin{aligned}
- E_{x \in X} \left[ \frac{\partial_i \partial_j p(x|\theta)}{p(x|\theta)} \right] &= \sum_{x \in X} [\partial_i \partial_j p(x|\theta)] \\
&= \partial_i \partial_j \sum_{x \in X} [p(x|\theta)] \\
&= - \partial_i \partial_j E_{x \in X} [1] \\
&= 0.
\end{aligned} \tag{A.2.3}$$

Here we assumed that we can swap the sum and the derivatives. For continuous variables, the swapping of the integral and the derivatives has to be assumed for the equation to hold.